

Boosting Algorithms

서울대학교 컴퓨터공학부 이영기



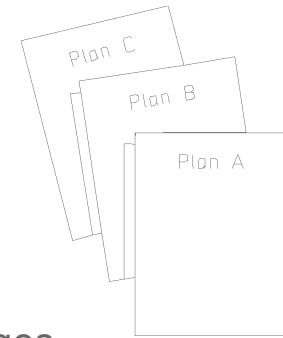
Human-Centered
Computer Systems Lab



SEOUL NATIONAL UNIVERSITY

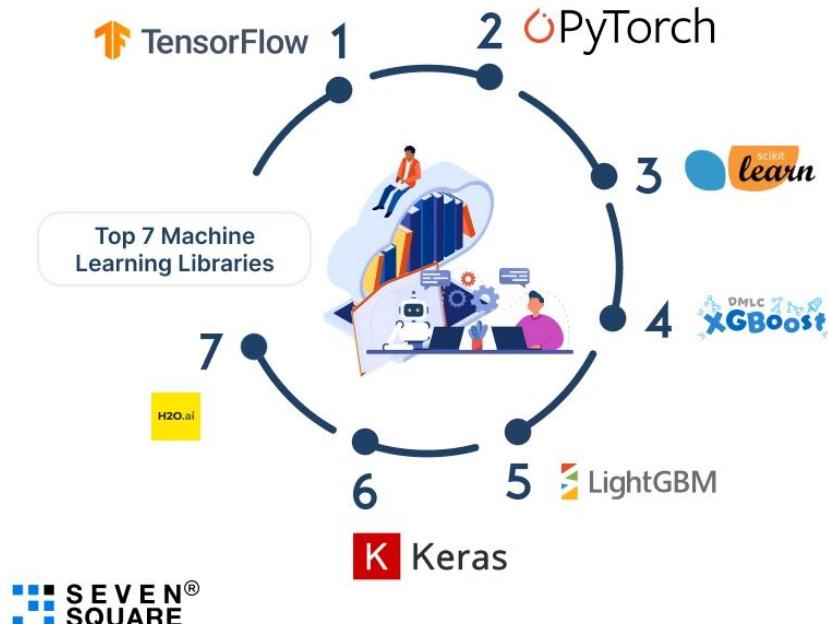
Overview

- **Objective**
 - Understand boosting algorithms
- **Content**
 - GBM(Gradient Boosting Machine)
 - XGBoost(eXtreme Gradient Boost)
 - LightGBM(Light Gradient-Boosting Machine)
- **After this module, you should be able to**
 - Understand various types of boosting and their advantages and disadvantages



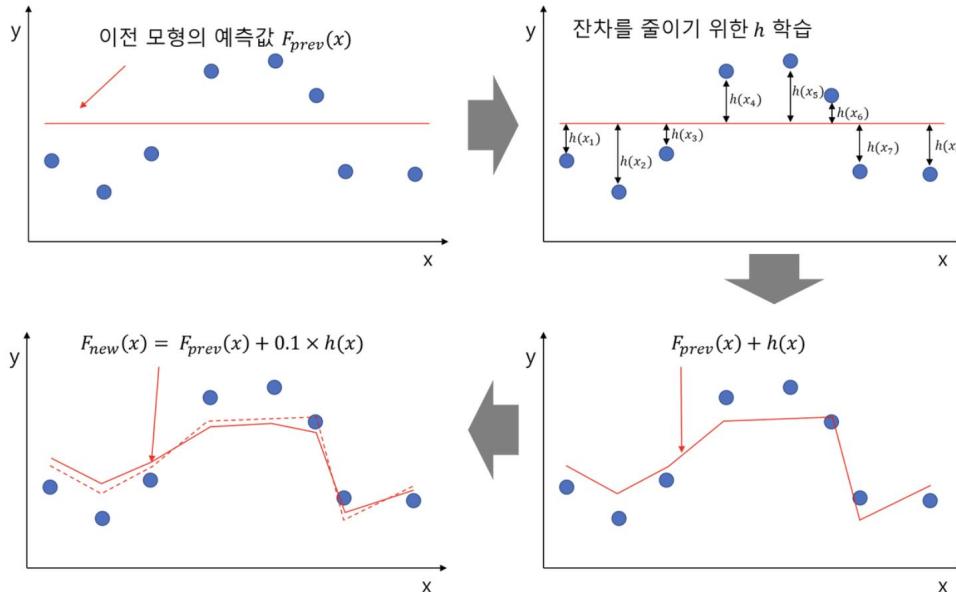
Introduction

Top 7 Machine Learning Libraries That You Should Explore



GBM(Gradient Boosting Machine)

- A **sequential** ensemble technique that builds strong predictive models by iteratively adding weak learners to correct the *residuals* of previous steps



$$F_m(x) = h_0(x) + l h_1(x) + \cdots + l h_m(x)$$

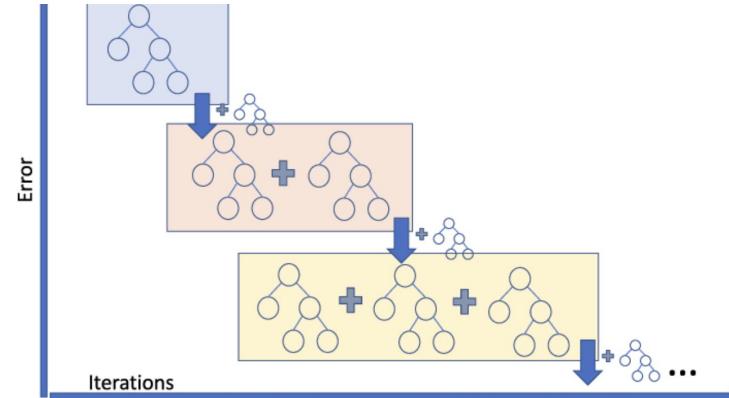
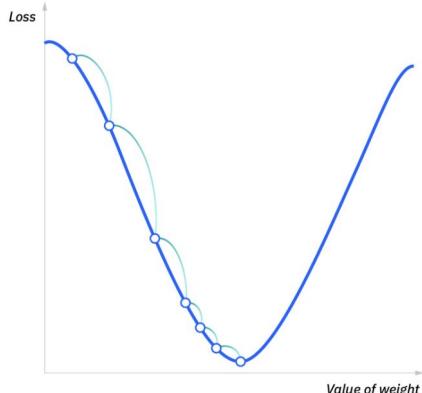
GBM(Gradient Boosting Machine)

(-) *residual*

- Gradient serves as a mathematical generalization of the *residual*

$$L(y_i, f(x_i)) = \frac{1}{2}(y_i - F(x_i))^2 \quad \frac{\partial L(y_i, F(x_i))}{\partial F(x_i)} = \frac{\partial[\frac{1}{2}(y_i - F(x_i))^2]}{\partial F(x_i)} = -(y_i - F(x_i))$$

- It performs **Gradient Descent**, following the steepest descent direction to minimize the overall loss function
- However, it is slow(working sequentially) and vulnerable to overfitting



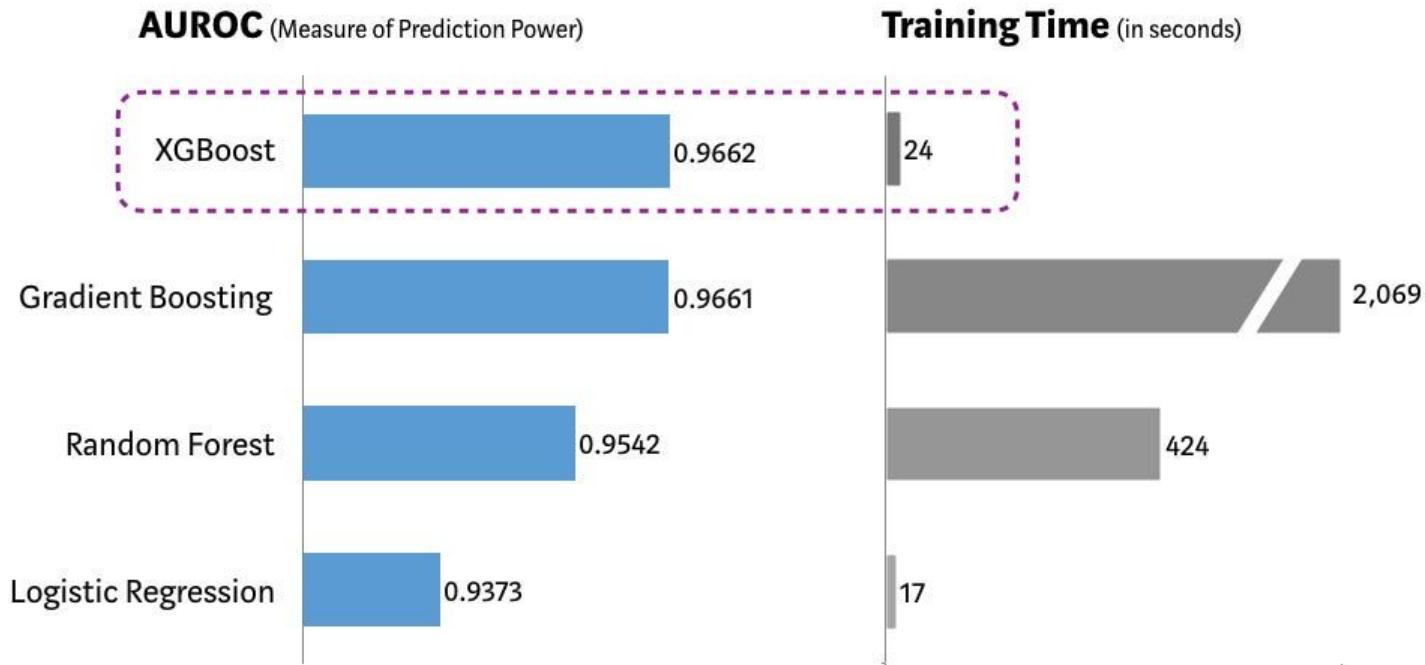
XGBoost(eXtreme Gradient Boost)

- A highly **scalable** and **regularized** gradient boosting framework to maximize computational speed and predictive accuracy through 2nd-order optimization
- Advantages
 - **Accuracy:** Uses 2nd-order Taylor Expansion (Hessian) for more accurate convergence compared to GBM's 1st-order gradients
 - **Regularization:** Prevents the model from becoming overfitting by using L2-Regularization
 - **Speed:** Maximizes speed by parallelizing split finding via pre-sorted column blocks and hardware-aware optimizations

XGBoost(eXtreme Gradient Boost)

Performance Comparison using SKLearn's 'Make_Classification' Dataset

(5 Fold Cross Validation, 1MM randomly generated data sample, 20 features)



CART(Classification And Regression Tree) Model

- A **binary decision tree** that assigns a real-valued score (w) to each leaf node, regardless of whether the task is classification or regression
- Key Attributes: Binary Splitting and Numerical Leaf Scores
- Essential for **Additive Learning**: These numerical scores act as correction values that are summed across all trees to form the final ensemble prediction

$$Y' = a \times \text{tree A} + b \times \text{tree B} + c \times \text{tree C} + \dots$$

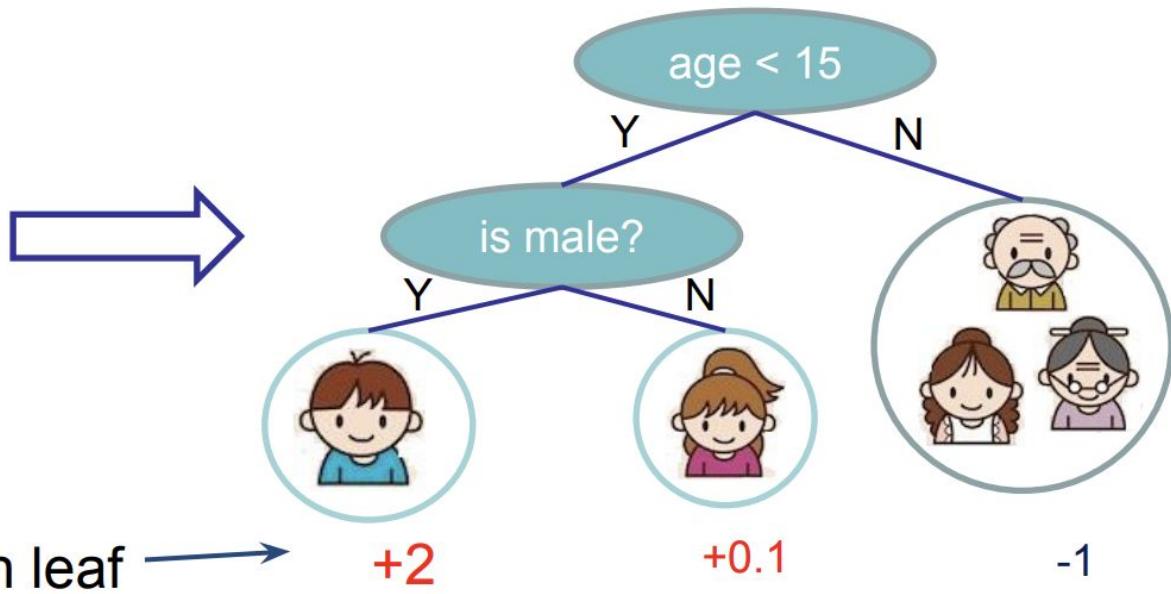
where a, b, c, \dots = weights for each tree

CART(Classification And Regression Tree) Model

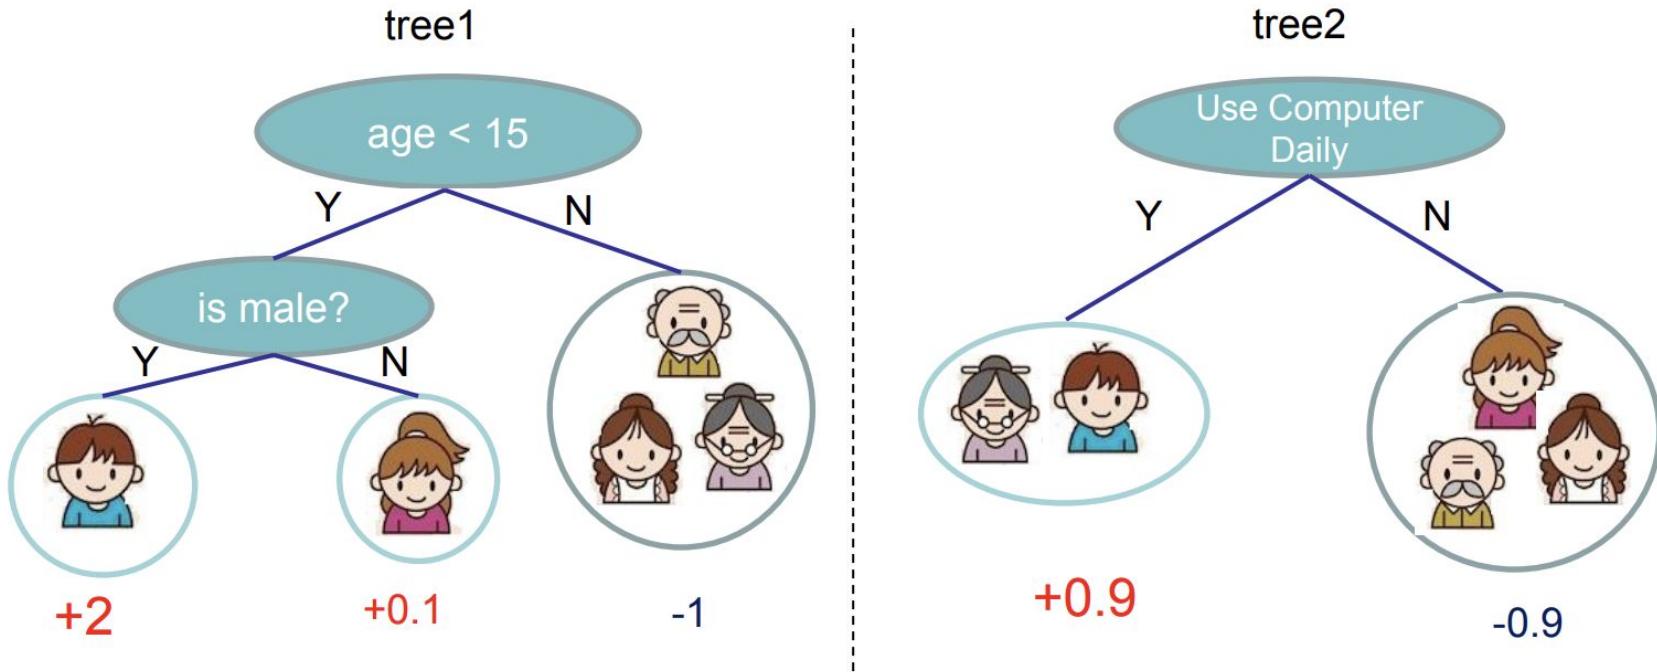
Input: age, gender, occupation, ...



Does the person like computer games



CART(Classification And Regression Tree) Model



$$f(\text{boy}) = 2 + 0.9 = 2.9$$

$$f(\text{old man}) = -1 - 0.9 = -1.9$$

XGBoost in a Nutshell

- Input: $n \times m$ (n samples, m features) $X = \{(\mathbf{x}_i, y_i)\}$, ($|X| = n, \mathbf{x}_i \in \mathbb{R}^m, y_i \in \mathbb{R}$)
- Model: \mathcal{F} (consists of B CART trees)

$$\hat{y}_i = \phi(\mathbf{x}_i) = \sum_{b=1}^B f_b(\mathbf{x}_i), f_b \in \mathcal{F}$$

where $\mathcal{F} = \{f(\mathbf{x}) = w_{q(\mathbf{x})}\}$ ($q : \mathbb{R}^m \rightarrow |T|, w \in \mathbb{R}^{|T|}$)

- Loss Term

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_b \Omega(f_b), \text{ where } \Omega(f) = \gamma|T| + \frac{1}{2}\lambda\|w\|^2$$

- Training

- Incremental branch expansion: Each tree grows by iteratively adding one branch at a time
- Post-pruning effect: By evaluating the Gain against γ , XGBoost effectively prunes unnecessary branches that do not contribute significantly to the model's predictive power

XGBoost: Regularization

- In XGBoost loss, we can see the regularization term (Ω)

$$\mathcal{L}(\phi) = \sum_i l(\hat{y}_i, y_i) + \sum_b \Omega(f_b), \text{ where } \Omega(f) = \gamma|T| + \frac{1}{2}\lambda\|w\|^2$$

- Unlike traditional GBM, regularization penalizes the number of leaves ($|T|$) and the magnitude of leaf weights (w , represented as L2-norm) to prevent overfitting
- Setting γ and λ to 0 reverts XGBoost to the standard Gradient Tree Boosting

XGBoost: Algorithm Details

Goal: find f_b to minimize loss($\mathcal{L}^{(b)}$)

- Loss Function

$$\mathcal{L}^{(b)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(b-1)} + f_b(\mathbf{x}_i)) + \Omega(f_b) + \text{Constant}$$

- It is easy to express loss when it is MSE, but what if we want to use other loss?
 - (i) Function can be complicated, (ii) Library should recalculate when loss term changes
- In XGBoost, we use **second-order approximation** by Taylor Expansion

$$\begin{aligned}\mathcal{L}^{(b)} &\simeq \sum_{i=1}^n [l(y_i, \hat{y}_i^{(b-1)}) + g_i f_b(\mathbf{x}_i) + \frac{1}{2} h_i f_b^2(\mathbf{x}_i)] + \Omega(f_t) \\ g_i &= \frac{\partial}{\partial \hat{y}_i^{(b-1)}} l(y_i, \hat{y}_i^{(b-1)}), \quad h_i = \frac{\partial^2}{\partial (\hat{y}_i^{(b-1)})^2} l(y_i, \hat{y}_i^{(b-1)})\end{aligned}$$

- Optimization depends only g_i & h_i , regardless of loss type, achieving better accuracy using Hessian (Newton's method)

XGBoost: Algorithm Details

- By applying the regularization term(Ω),

$$\tilde{\mathcal{L}}^{(b)} = \sum_{i=1}^n [g_i f_b(x_i) + \frac{1}{2} h_i f_b^2(x_i)] + \gamma |T| + \frac{1}{2} \lambda \sum_{j=1}^{|T|} w_j^2$$

- Using a new set $I_j = \{i | q(x_i) = j\}$

$$\tilde{\mathcal{L}}^{(b)} = \sum_{j=1}^{|T|} \left[\left(\sum_{i \in I_j} g_i \right) w_j + \frac{1}{2} \left(\sum_{i \in I_j} h_i + \lambda \right) w_j^2 \right] + \gamma |T|$$

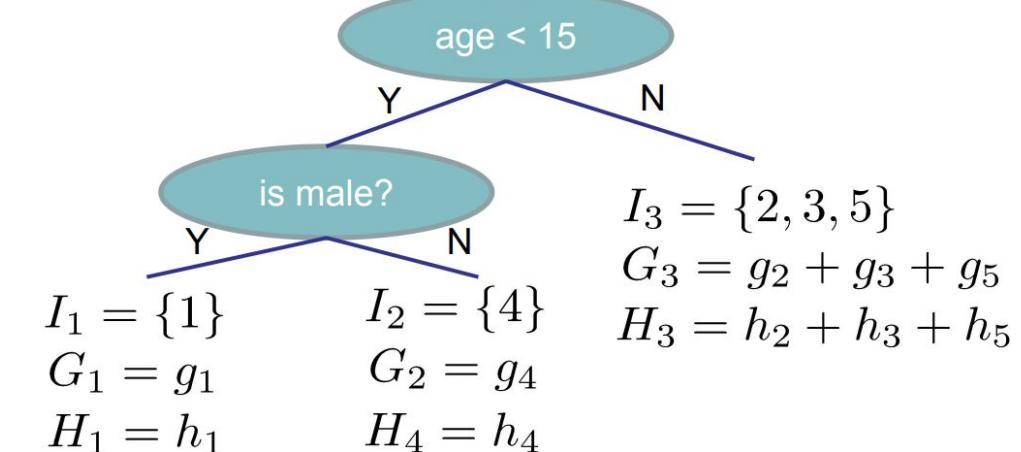
- By differentiating the objective function with respect to w_j , we obtain the optimal weight w_j^* that minimizes the loss for each leaf, and calculate the optimized loss

$$w_j^* = -\frac{\sum_{i \in I_j} g_i}{\sum_{i \in I_j} h_i + \lambda} \quad \tilde{\mathcal{L}}^{(b)}(q) = -\frac{1}{2} \sum_{j=1}^{|T|} \frac{\left(\sum_{i \in I_j} g_i \right)^2}{\sum_{i \in I_j} h_i + \lambda} + \gamma |T|$$

XGBoost: Example

Instance index gradient statistics

1		g1, h1
2		g2, h2
3		g3, h3
4		g4, h4
5		g5, h5



$$Obj = - \sum_j \frac{G_j^2}{H_j + \lambda} + 3\gamma$$

The smaller the score is, the better the structure is

XGBoost: Greedy Update

- Evaluating the gain of a potential split

- Compare the loss(org - split) when split happens: $I = I_L \cup I_R$

$$\mathcal{L}_{split} = \frac{1}{2} \left[\frac{(\sum_{i \in I_L} g_i)^2}{\sum_{i \in I_L} h_i + \lambda} + \frac{(\sum_{i \in I_R} g_i)^2}{\sum_{i \in I_R} h_i + \lambda} - \frac{(\sum_{i \in I} g_i)^2}{\sum_{i \in I} h_i + \lambda} \right] - \gamma$$

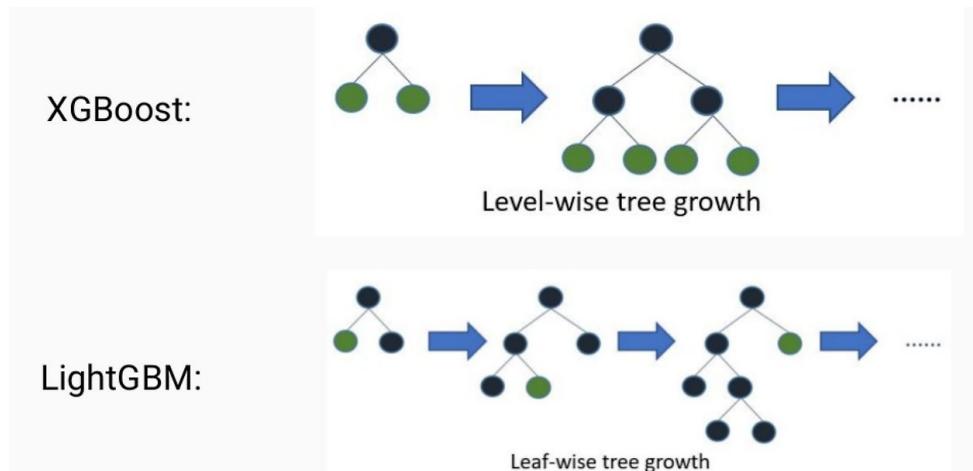
- If $\mathcal{L}_{split} > 0$, splitting benefits the optimization -> Keep splitting
- γ, λ : regularization parameter
 - If they get larger, $\mathcal{L}_{split} < 0$
 - Prevents overfitting(complex tree)

LightGBM: Overcoming the Scalability Bottleneck

- The most time-consuming part of GBM is **finding the optimal split point**, which typically requires scanning every data point (N) for every feature (D)
 - $O(N \times D)$ complexity becomes a massive burden for massive dataset
- Reducing N : GOSS (Gradient-based One-Side Sampling)
 - Prioritize high-gradient instances as they represent **under-fitted data** with the most information gain
- Reducing D : EFB (Exclusive Feature Bundling)
 - **Bundle mutually exclusive** sparse features(rarely non-zero simultaneously) into single dense feature

LightGBM: Leaf-wise Tree Growth

- LightGBM grows the tree by **splitting the leaf node** that results in the greatest loss reduction(max gain), regardless of its current depth
 - Efficient(achieve lower loss faster), but forms asymmetric tree(overfitting)
- Tuning parameters
 - *num_leaves, max_depth, min_data_in_leaf*



LightGBM vs. XGBoost

- LightGBM achieves target accuracy **much faster** than other GBM algorithms

