

LG전자 고급 데이터 사이언티스트 양성 과정 7기

Automated Machine Learning (AutoML)

서울대학교 컴퓨터공학부 이영기

Human-Centered Computer Systems Lab




SEOUL NATIONAL UNIVERSITY

Overview


- **Objective**
 - Identifying challenges in traditional ML pipeline
 - Learn various Hyperparameter Optimization (HPO) techniques
- **Content**
 - Introduction to AutoML
 - Hyperparameter Optimization
 - Blackbox Optimization(Grid search, Random search, Bayesian optimization, Tree-structured Parzen Estimator)
 - Multi-fidelity Optimization (Successive halving, Hyperband, Bayesian optimization + Hyperband)

Traditional Machine Learning Pipeline


- Building an ML model is an iterative, complex, and time-consuming process
 - Clean & preprocess the data
 - Select / engineer better features
 - Select an algorithm
 - Set the model parameters
 - ...



Challenges in Designing ML Pipelines



Complex
Search Space




Black-Box
Problem




Expensive
Evaluations




Noise on
observations




From Manual ML to Automated ML



Design Decisions taken care by AutoML



Hyper-
parameters



Architecture
Design

preprocessor	#λ
extreml. rand. trees prepr.	5
fast ICA	4
feature agglomeration	4
kernel PCA	5
rand. kitchen sinks	2
linear SVM prepr.	3
no preprocessing	-
nystroem sampler	5
PCA	2
polynomial	3
random trees embed.	4
select percentile	2
select rates	3
<hr/>	
one-hot encoding	2
imputation	1
balancing	1
rescaling	1


Pre-
processing

classifier	#λ
AdaBoost (AB)	4
Bernoulli naïve Bayes	2
decision tree (DT)	4
extreml. rand. trees	5
Gaussian naïve Bayes	-
gradient boosting (GB)	6
kNN	3
LDA	4
linear SVM	4
kernel SVM	7
multinomial naïve Bayes	2
passive aggressive	3
QDA	2
random forest (RF)	5
Linear Class. (SGD)	10

Algorithms

Importance of Design Decisions in ML


- Each decision is critical to overall performance
 - E.g. Classifiers performance compared to auto-sklearn over time in one OpenML dataset



Importance of Design Decisions in ML

- Each decision is critical to overall performance
 - E.g. Classifiers performance compared to auto-sklearn over time in one OpenML dataset

Choosing the gradient boosting algorithm instead of kernel SVM
→ 17% improv.



Optimizing hyperparameters in gradient boosting
→ 20% improv.

Hyperparameter Optimization (HPO)

Types of Hyperparameters

- **Continuous**
 - Regularization parameter C in SVM
 - Shrinkage parameter λ in boosting
- **Integer**
 - Number of trees in boosting and random forest
- **Categorical**
 - Algorithm (SVM, RF, NN, ...)
 - SVM kernel (Polynomial, RBF, ...)
 - Activation function (ReLU, Leaky ReLU, ...)

Types of Hyperparameters

- Conditional: hyperparameters B are only active if other hyperparameters A are set a certain way
 - Example 1:
 - A = choice of optimizer (Adam or SGD)
 - B = Adam's second momentum hyperparameter (only active if A=Adam)
 - Example 2:
 - A = type of layer k (convolution, max pooling, fully connected, ...)
 - B = conv. kernel size of that layer (only active if A = convolution)
 - Example 3:
 - A = choice of classifier (RF or SVM)
 - B = SVM's kernel parameter (only active if A = SVM)

AutoML as Hyperparameter Optimization

Definition: Combined Algorithm Selection and Hyperparameter Optimization (CASH)

Let

- $\mathcal{A} = \{A^{(1)}, \dots, A^{(n)}\}$ be a set of algorithms
- $\Lambda^{(i)}$ denote the hyperparameter space of $A^{(i)}$, for $i = 1, \dots, n$
- $\mathcal{L}(A_{\lambda}^{(i)}, D_{train}, D_{valid})$ denote the loss of $A^{(i)}$, using $\lambda \in \Lambda^{(i)}$ trained on D_{train} and evaluated on D_{valid} .

The Combined Algorithm Selection and Hyperparameter Optimization (CASH) problem is to find a combination of algorithm $A^* = A^{(i)}$ and hyperparameter configuration $\lambda^* \in \Lambda^{(i)}$ that minimizes this loss:

$$A_{\lambda^*}^* \in \arg \min_{A^{(i)} \in \mathcal{A}, \lambda \in \Lambda^{(i)}} \mathcal{L}(A_{\lambda}^{(i)}, D_{train}, D_{valid})$$


→ Simply a HPO problem with a top-level hyperparameter (choice of algorithm) that all other hyperparameters are conditional on.

Hyperparameter Optimization

- Blackbox optimization
 - Grid search
 - Random search
 - Bayesian optimization
 - Tree-structured Parzen Estimator
- Multi-fidelity optimization
 - Successive Halving
 - Hyperband
 - Bayesian optimization + Hyperband


Blackbox Hyperparameter Optimization

- The blackbox function is expensive to evaluate
→ Can't sample too often!





Blackbox Optimization: Grid and Random Search

- Grid search: required number of samples grows exponentially
- Random search: handles unimportant dimensions better



Blackbox Optimization: Grid and Random Search

- Grid search: required number of samples grows exponentially
- Random search: handles unimportant dimensions better




Bayesian Optimization

- Limitation of previous techniques: completely uninformed and does not learn about the search space
- If we have already evaluated a set of candidates, can't we do better than random selection?

=> Predict promising regions based on history of observations


Bayesian Optimization

- Update the surrogate model sequentially with new observations
- Use the acquisition function to select new promising configurations to evaluate
- Surrogate Model
 - Probabilistic modeling of $L(\lambda) \sim (\hat{L}(\lambda), \hat{\sigma}(\lambda))$ with posterior mean $\hat{L}(\lambda)$ and uncertainty $\hat{\sigma}(\lambda)$
E.g. Gaussian Processes
- Acquisition Function
 - Balance exploration (high $\hat{\sigma}$) vs. exploitation (low \hat{L})
 - E.g. Expected improvement (EI):
$$u(\lambda) = E[\max\{L_{\min} - L(\lambda), 0\}],$$
 where L_{\min} is lowest value from observations




Bayesian Optimization

1. Start with a few random hyperparameter configurations
2. Fit a surrogate model to observations
3. Compute the acquisition function and select new configurations accordingly
4. Evaluate new points and get observations
5. Repeat 2-4




Bayesian Optimization

1. Start with a few random hyperparameter configurations
2. Fit a surrogate model to observations
3. Compute the acquisition function and select new configurations accordingly
4. Evaluate new points and get observations
5. Repeat 2-4




Bayesian Optimization

1. Start with a few random hyperparameter configurations
2. Fit a surrogate model to observations
3. Compute the acquisition function and select new configurations accordingly
4. Evaluate new points and get observations
5. Repeat 2-4




Bayesian Optimization



Tree-structured Parzen Estimator (TPE)

- Problems for standard Gaussian Process approach:
 - Not scalable to high dimensional hyperparameter spaces
 - Unfit for discrete spaces since GP assumes smoothness (e.g. Conditional hyperparameters)
- Maintain two surrogate models
 $p(\lambda \text{ is good})$ and $p(\lambda \text{ is bad})$,
rather than $p(L | \lambda)$
 - $P(\lambda \text{ is good}) = P(\lambda | L \leq L^*)$
 - $P(\lambda \text{ is bad}) = P(\lambda | L > L^*)$
(L^* : split threshold)
- Acquisition function $\propto \frac{P(\lambda \text{ is good})}{P(\lambda \text{ is bad})}$
 - Equivalent to expected improvement




Hyperparameter Optimization

- Blackbox optimization
 - Grid search
 - Random search
 - Bayesian optimization
 - Tree-structured Parzen Estimator
- Multi-fidelity optimization
 - Successive Halving
 - Hyperband
 - Bayesian optimization + Hyperband

Multi-fidelity Optimization

- Massive dataset sizes and complex models make blackbox performance evaluation expensive




Multi-fidelity Optimization

- Massive dataset sizes and complex models make blackbox performance evaluation expensive
- Make use of cheap low-fidelity evaluations with a small budget, e.g.
 - Subsets of the data
 - Fewer epochs of iterative training algorithms
- These approximations introduce a tradeoff between optimization performance and runtime


Successive Halving

- Many cheap evaluations on small subsets & few expensive evaluations on the full data
- E.g. cheap low-fidelity evaluation of SVM, with subsets of the MNIST data



Successive Halving


- Many cheap evaluations on small subsets & few expensive evaluations on the full data
- E.g. cheap low-fidelity evaluation of SVM, with subsets of the MNIST data



- Up to 1000x speedups

Successive Halving

- Successive halving for eight algorithms/configurations
- After evaluating all algorithms on 1/8 of the total budget, half of them are dropped and the budget given to the remaining algorithms is doubled







Successive Halving: Limitation

- Budget vs. number tradeoff in SHA
 - Given a total budget, the user has to decide whether:
 - Try many configurations and only assign a small budget to each
 - Try only a few and assign them a larger budget
- Assigning too small budget can result in prematurely terminating good configurations
- Assigning too large budget can result in running poor configurations too long

HyperBand

- Trying various budget-number combinations
 - Repeat SHA with different initial number of configurations n and start budgets m
 - Randomly select the initial configurations for each iteration
 - Return configuration with the smallest intermediate loss seen so far



High n & Low m

Low n & High m


BOHB: Bayesian Optimization and HyperBand

- Advantage of Bayesian optimization
 - Strong final performance
- Advantage of HyperBand
 - Strong anytime performance
 - Easily parallelizable
 - Scalable

BOHB: Bayesian Optimization and HyperBand


- Advantage of Bayesian optimization
 - Strong final performance
- Advantage of HyperBand
 - Strong anytime performance
 - Easily parallelizable
 - Scalable
- Combining the best of both: BOHB
 - Bayesian optimization (BO): **choosing the configuration** to evaluate (using TPE)
 - HyperBand (HB): deciding **how to allocate budgets**

HyperBand vs Random Search




Biggest advantage: much improved **anytime performance**

Bayesian Optimization vs Random Search



Biggest advantage: much improved **final performance**

Combining Bayesian Optimization & HyperBand



Best of both worlds: strong **anytime and final performance**

BOHB Algorithm

- Relies on HB to determine number of configurations and budget to evaluate
- Replaces the random selection of configurations at the beginning of each HB iteration by a TPE model-based search
- Update TPE using intermediate losses during SHA, prioritize results from higher budget
 - Use budget $b = \text{argmax} \{ D_b : |D_b| \geq N_{min} + 2\}$, where D_b is observation with budget b and N_{min} is minimum number of points to build a model

References

- AutoML tutorial at NeurIPS 2018¹⁾
- AutoML-101²⁾
- Scalable Machine Learning and Deep Learning, KTH³⁾
- Automated machine learning⁴⁾

1) <https://www.youtube.com/watch?v=0eBR8a4MQ30>

2) <https://www.automl.org/wp-content/uploads/2021/03/AutoML-101.pdf>

3) https://id2223kth.github.io/slides/2022/11_automl.pdf

4) Hutter, Frank, Lars Kotthoff, and Joaquin Vanschoren. *Automated machine learning: methods, systems, challenges*. Springer Nature, 2019.